

College Data Pipeline

Betsy Fridman, Alex Hill, Ellie Huang, Steve Wong

Project Overview

Build a unified, year-over-year relational database combining **College Scorecard** and **IPEDS** data to enable fast, historical queries on U.S. higher-education institutions.

Main Objectives

- Build a longitudinal SQL database integrating annual College Scorecard and IPEDS data.
- Link datasets using UNITID/OPEID crosswalks.
- Track year-over-year changes in admissions, tuition, outcomes, and institutional characteristics.
- Store geographic and Carnegie Classification info for historical analysis.
- Enable fast lookups and trend analyses for any institution and year.

About the Data

General Institution Information (location, institution type)

- Source: National Center for Education Statistics - Integrated Postsecondary Education Data System (IPEDS)
- Documentation: Click “Complete Data Files”, and then the year
- Upload frequency: 3 times a year (Fall, Winter, Spring), last update was on September 23, 2025

Scorecard Data: More Detailed Institution Information (financials, earnings)

- Source: U.S. Department of Education (College Scorecard)
- Documentation
- Upload frequency: Annually, last update was on November 17, 2025

Data to Store: IPEDS

- Institution ID (unique ID, primary key)
- Institution name
- Institution location data (address, city, state abbreviation, zip etc.)
- Carnegie classification
- Year (the end year for that academic period)

Data to Store: College Scorecard

- Institution ID (unique ID)
- Accreditation agency (grants formal recognition to educational institutions)
- Degree (predominant, highest)
- Control (private/public/nonprofit)
- Region
- Admission rate
- Tuition information (in-state, out-of-state, or program-specific, also revenue from tuition per student)
- Average faculty salary
- Default rates


Schema Overview

1 main table

- **institution_ipeds_info** (parent table, every other table references this one)
→ institution-specific information that doesn't change very often

4 reference tables

- **institution_scorecard_info** → more institution info (Updated annually)
- **institution_financials** → yearly financial metrics
- **institution_admissions** → varies year to year, testing metrics
- **institution_completion** → graduation rates and demographic breakdowns/annual student outcomes



Set-up

1. Clone [GitHub repository](#)
2. Create local files (credentials_copy.py to store credentials)
3. Store data in root folder

Recommended Repo Structure

```
└── College_Scorecard_export (data zip  
    files)  
└── README.md  
└── credentials_copy.py  
└── part_two.ipynb  
└── load-ipeds.py  
└── load-scorecard.py
```

Execute Order:

1. Create `credentials_copy.py` + upload sensitive user/database information
2. `part_two.ipynb` to create tables
3. [load-ipeds.py](#) FIRST (four times)
4. Then, [load-scorecard.py](#). (four times as well)

Preprocessing IPEDS data

Column Selection & Validation: Selects columns (15 variables) format and creates a year variable, added as a new column.

Date & Time Handling: (To do) need to implement to ensure consistency across time zones, formatting, etc.

Numeric Data Cleaning: Replaces NA and nan with None across all 16 metric columns.

Preprocessing Scorecard data

Column Selection & Filtering: Selects required columns from scorecard data, then filters to only valid Unit IDs (valid = matches in IPEDS dataset, otherwise drop the row)

Ensure UNITID exists in scorecard data first.

Load script - IPEDS

Load CSV Input

- Takes file path from `sys.argv[1]`
- Reads the CSV into a pandas DataFrame
- Extracts the year from the filename and adds it as a `YEAR` column

Database Connection

- Uses credentials from `credentials_copy.py`
- Connects to the Postgres database with `psycopg2`

Insert/Update Logic (`insert_dataframe()`)

- Keeps only the columns needed for the target table
- For each row in the DataFrame:
 - If `UNITID` exists in the database: **update** the existing row
 - If `UNITID` does not exist: **insert** a new row
- Processes the data row-by-row until complete

Load script - Scorecard

Load and Prepare Data

- Takes CSV path from `sys.argv[1]` and loads it with `pd.read_csv()`
- Extracts YEAR from filename and adds **YEAR + 1** (Scorecard release lag)
- Splits the large Scorecard file into four cleaned dataframes

Database Prep

- Connects using credentials from `credentials_copy.py`
- Queries all valid UNITIDs from `institution_ipeds_info`
- Filters each Scorecard dataframe to only keep UNITIDs that already exist
- Prints counts of kept vs dropped rows

Insert Process (`insert_dataframe_strict()`)

- Takes a cleaned dataframe, target table name, and DB credentials, connect to database.
- Inserts rows **one by one**
 - If **any row fails**, the function **stops immediately** and performs a full **ROLLBACK**
 - If all rows succeed, performs a single **COMMIT**

Batching Strategy - (To do: look into faster processing)

- **Performance Optimization:** Executing 500 rows at once with `executemany()` is significantly faster than 500 individual `execute()` calls, reducing database round trips and network overhead
- **Transaction Efficiency:** Each batch commits as a unit, balancing between performance (larger batches) and recovery speed (smaller batches when errors occur)

Error Handling

- If a batch fails, print the error
 - Start inserting the rows in the batch one-by-one
 - Stop when the problematic row is found
 - Print the exact row and the reason for the error
 - Stop execution of program
- No data is inserted when an error occurs